

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv(r'/Users/divyanshsahai/Desktop/breast_cancer.csv')
df = df.loc[:, ~df.columns.str.contains('^Unnamed')]
```

```
In [5]: df.head()
```

```
Out[5]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactnes
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	

5 rows × 32 columns

```
In [6]: df.shape
```

```
Out[6]: (569, 32)
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactnes
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000	569
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360	0
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014064	0
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630	0
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370	0
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870	0
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300	0
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400	0

8 rows × 31 columns

```
In [8]: df.diagnosis.unique()
```


```
Out[8]: array(['M', 'B'], dtype=object)
```

```
In [9]: df['diagnosis'].value_counts()
```

```
Out[9]: diagnosis
B      357
M      212
Name: count, dtype: int64
```

```
In [22]: sns.countplot(x='diagnosis', data=df, palette='hls')
plt.savefig('Logistic.png')
```

```
plt.show()
```



```
In [12]: df.drop('id',axis=1,inplace=True)
```

```
In [13]: df.head()
```

```
Out[13]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	M	17.99	10.38	122.80	1001.0	0.11840	0.27760
1	M	20.57	17.77	132.90	1326.0	0.08474	0.07864
2	M	19.69	21.25	130.00	1203.0	0.10960	0.15990
3	M	11.42	20.38	77.58	386.1	0.14250	0.28390
4	M	20.29	14.34	135.10	1297.0	0.10030	0.13280

5 rows × 31 columns

```
In [14]: df['diagnosis'] = df['diagnosis'].map({'M':1,'B':0})
df.head()
```

```
Out[14]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	1	17.99	10.38	122.80	1001.0	0.11840	0.27760
1	1	20.57	17.77	132.90	1326.0	0.08474	0.07864
2	1	19.69	21.25	130.00	1203.0	0.10960	0.15990
3	1	11.42	20.38	77.58	386.1	0.14250	0.28390
4	1	20.29	14.34	135.10	1297.0	0.10030	0.13280

5 rows × 31 columns

```
In [15]: df.isnull().sum()
```

```
Out[15]: diagnosis                0
radius_mean                    0
texture_mean                   0
perimeter_mean                 0
area_mean                     0
smoothness_mean                0
compactness_mean               0
concavity_mean                 0
concave points_mean            0
symmetry_mean                  0
fractal_dimension_mean         0
radius_se                      0
texture_se                     0
perimeter_se                   0
area_se                        0
smoothness_se                  0
compactness_se                 0
concavity_se                   0
concave points_se              0
symmetry_se                    0
fractal_dimension_se           0
radius_worst                   0
texture_worst                  0
perimeter_worst                0
area_worst                     0
smoothness_worst               0
compactness_worst              0
concavity_worst                0
concave points_worst           0
symmetry_worst                 0
fractal_dimension_worst        0
dtype: int64
```


```
In [16]: df.corr()
```

Out[16]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
diagnosis	1.000000	0.730029	0.415185	0.742636	0.708984	0.358560
radius_mean	0.730029	1.000000	0.323782	0.997855	0.987357	0.170581
texture_mean	0.415185	0.323782	1.000000	0.329533	0.321086	-0.023389
perimeter_mean	0.742636	0.997855	0.329533	1.000000	0.986507	0.207278
area_mean	0.708984	0.987357	0.321086	0.986507	1.000000	0.177028
smoothness_mean	0.358560	0.170581	-0.023389	0.207278	0.177028	1.000000
compactness_mean	0.596534	0.506124	0.236702	0.556936	0.498502	0.659123
concavity_mean	0.696360	0.676764	0.302418	0.716136	0.685983	0.521984
concave points_mean	0.776614	0.822529	0.293464	0.850977	0.823269	0.553695
symmetry_mean	0.330499	0.147741	0.071401	0.183027	0.151293	0.557775
fractal_dimension_mean	-0.012838	-0.311631	-0.076437	-0.261477	-0.283110	0.584792
radius_se	0.567134	0.679090	0.275869	0.691765	0.732562	0.301467
texture_se	-0.008303	-0.097317	0.386358	-0.086761	-0.066280	0.068406
perimeter_se	0.556141	0.674172	0.281673	0.693135	0.726628	0.296092
area_se	0.548236	0.735864	0.259845	0.744983	0.800086	0.246552
smoothness_se	-0.067016	-0.222600	0.006614	-0.202694	-0.166777	0.332375
compactness_se	0.292999	0.206000	0.191975	0.250744	0.212583	0.318943
concavity_se	0.253730	0.194204	0.143293	0.228082	0.207660	0.248396
concave points_se	0.408042	0.376169	0.163851	0.407217	0.372320	0.380676
symmetry_se	-0.006522	-0.104321	0.009127	-0.081629	-0.072497	0.200774
fractal_dimension_se	0.077972	-0.042641	0.054458	-0.005523	-0.019887	0.283607
radius_worst	0.776454	0.969539	0.352573	0.969476	0.962746	0.213120
texture_worst	0.456903	0.297008	0.912045	0.303038	0.287489	0.036072
perimeter_worst	0.782914	0.965137	0.358040	0.970387	0.959120	0.238853
area_worst	0.733825	0.941082	0.343546	0.941550	0.959213	0.206718
smoothness_worst	0.421465	0.119616	0.077503	0.150549	0.123523	0.805324
compactness_worst	0.590998	0.413463	0.277830	0.455774	0.390410	0.472468
concavity_worst	0.659610	0.526911	0.301025	0.563879	0.512606	0.434926
concave points_worst	0.793566	0.744214	0.295316	0.771241	0.722017	0.503053
symmetry_worst	0.416294	0.163953	0.105008	0.189115	0.143570	0.394309
fractal_dimension_worst	0.323872	0.007066	0.119205	0.051019	0.003738	0.499316

31 rows × 31 columns

```
In [17]: plt.figure(figsize=(20,20))
sns.heatmap(df.corr(), annot=True)
plt.savefig('CorrelationMatrix.png')
```



```
In [21]: # Generate and visualize the correlation matrix
corr = df.corr().round(2)


# Mask for the upper triangle
mask = np.zeros_like(corr, dtype=bool)
mask[np.triu_indices_from(mask)] = True

# Set figure size
f, ax = plt.subplots(figsize=(20, 20))

# Define custom colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap
sns.heatmap(corr, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5}, annot=True)
```

```
plt.tight_layout()
plt.savefig('Heatmap.png')
```



```
In [23]: # first, drop all "worst" columns
cols = ['radius_worst',
        'texture_worst',
        'perimeter_worst',
        'area_worst',
        'smoothness_worst',
        'compactness_worst',
        'concavity_worst',
        'concave points_worst',
        'symmetry_worst',
        'fractal_dimension_worst']
df = df.drop(cols, axis=1)

# then, drop all columns related to the "perimeter" and "area" attributes
cols = ['perimeter_mean',
        'perimeter_se',
        'area_mean',
        'area_se']
df = df.drop(cols, axis=1)

# lastly, drop all columns related to the "concavity" and "concave points" attributes
cols = ['concavity_mean',
```

```

        'concavity_se',
        'concave points_mean',
        'concave points_se']
df = df.drop(cols, axis=1)

# verify remaining columns
df.columns

```

```

Out[23]: Index(['diagnosis', 'radius_mean', 'texture_mean', 'smoothness_mean',
        'compactness_mean', 'symmetry_mean', 'fractal_dimension_mean',
        'radius_se', 'texture_se', 'smoothness_se', 'compactness_se',
        'symmetry_se', 'fractal_dimension_se'],
        dtype='object')


```

```

In [25]: # Draw the heatmap again, with the new correlation matrix
corr = df.corr().round(2)
mask = np.zeros_like(corr, dtype=bool)
mask[np.triu_indices_from(mask)] = True

f, ax = plt.subplots(figsize=(20, 20))
sns.heatmap(corr, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5}, annot=True)
plt.tight_layout()
plt.savefig('RevisedHeatmap.png')

```



```
In [28]: X=df.drop(['diagnosis'],axis=1)
y = df['diagnosis']
```

```
In [29]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=40)
```

```
In [30]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()

X_train=ss.fit_transform(X_train)
X_test=ss.fit_transform(X_test)
```

```
In [31]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
```

```
In [32]: from sklearn.linear_model import LogisticRegression
```

```
In [33]: lr=LogisticRegression()

model1=lr.fit(X_train,y_train)
prediction1=model1.predict(X_test)
```

```
In [34]: from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]
lr_probs = model1.predict_proba(X_test)

# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]

# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)

# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))

# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)


# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')

# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')

# show the legend
pyplot.legend()

# show the plot
pyplot.show()
plt.savefig('ROCR.png')
```


No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.986




<Figure size 640x480 with 0 Axes>

```
In [35]: cm=confusion_matrix(y_test,prediction1)
cm
```

```
Out[35]: array([[110,  5],
               [ 2, 54]])
```

```
In [36]: sns.heatmap(cm,annot=True,cmap="YlOrBr",linecolor='White', linewidths=1)
plt.savefig('Logistic.png')
```



```
In [37]: accuracy_score(y_test, prediction1)
```

```
Out[37]: 0.9590643274853801
```

```
In [38]: print(classification_report(y_test, prediction1))
```

	precision	recall	f1-score	support
0	0.98	0.96	0.97	115
1	0.92	0.96	0.94	56
accuracy			0.96	171
macro avg	0.95	0.96	0.95	171
weighted avg	0.96	0.96	0.96	171

```
In [39]: from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from matplotlib import pyplot as plt
```

```
In [40]: dtc=DecisionTreeClassifier()
model2=dtc.fit(X_train,y_train)
prediction2=model2.predict(X_test)
```

```
In [41]: text_representation = tree.export_text(dtc)
print(text_representation)
```

```

|--- feature_0 <= 0.24
|   |--- feature_3 <= 0.40
|       |--- feature_1 <= 0.08
|           |--- feature_6 <= 0.72
|               |--- class: 0
|           |--- feature_6 > 0.72
|               |--- feature_0 <= -0.97
|                   |--- class: 0
|               |--- feature_0 > -0.97
|                   |--- class: 1
|       |--- feature_1 > 0.08
|           |--- feature_0 <= -0.22
|               |--- feature_2 <= 1.41
|                   |--- class: 0
|               |--- feature_2 > 1.41
|                   |--- feature_11 <= -0.29
|                       |--- class: 1
|                   |--- feature_11 > -0.29
|                       |--- class: 0
|           |--- feature_0 > -0.22
|               |--- feature_2 <= -0.28
|                   |--- feature_5 <= -0.87
|                       |--- feature_4 <= -0.93
|                           |--- class: 0
|                       |--- feature_4 > -0.93
|                           |--- class: 1
|                   |--- feature_5 > -0.87
|                       |--- class: 0
|               |--- feature_2 > -0.28
|                   |--- feature_5 <= 0.49
|                       |--- class: 1
|                   |--- feature_5 > 0.49
|                       |--- class: 0
|       |--- feature_3 > 0.40
|           |--- feature_0 <= -0.70
|               |--- class: 0
|           |--- feature_0 > -0.70
|               |--- feature_1 <= -0.94
|                   |--- class: 0
|               |--- feature_1 > -0.94
|                   |--- feature_4 <= 0.22
|                       |--- feature_3 <= 0.47
|                           |--- class: 1
|                       |--- feature_3 > 0.47
|                           |--- class: 0
|                   |--- feature_4 > 0.22
|                       |--- feature_1 <= -0.68
|                           |--- feature_2 <= 1.51
|                               |--- class: 0
|                           |--- feature_2 > 1.51
|                               |--- class: 1
|                       |--- feature_1 > -0.68
|                           |--- class: 1
|--- feature_0 > 0.24
|   |--- feature_1 <= -0.70
|       |--- feature_9 <= -0.16
|           |--- class: 0
|       |--- feature_9 > -0.16
|           |--- feature_0 <= 0.29
|               |--- class: 0
|           |--- feature_0 > 0.29
|               |--- class: 1
|   |--- feature_1 > -0.70
|       |--- feature_4 <= -1.14

```

```


| | | | |--- feature_11 <= -0.66
| | | | |--- class: 0
| | | | |--- feature_11 > -0.66
| | | | |--- class: 1
| | | |--- feature_4 > -1.14
| | | |--- class: 1

```

```
In [42]: cm2= confusion_matrix(y_test,prediction2)
cm2
```

```
Out[42]: array([[105, 10],
               [ 5, 51]])
```

```
In [43]: sns.heatmap(cm2,annot=True, cmap ="YlGn",linecolor='White', linewidths=1)
plt.savefig('DecisionTree.png')
```



```
In [44]: accuracy_score(y_test,prediction2)
```

```
Out[44]: 0.9122807017543859
```

```
In [45]: print(classification_report(y_test, prediction2))
```

	precision	recall	f1-score	support
0	0.95	0.91	0.93	115
1	0.84	0.91	0.87	56
accuracy			0.91	171
macro avg	0.90	0.91	0.90	171
weighted avg	0.92	0.91	0.91	171


```
In [46]: from sklearn.ensemble import RandomForestClassifier
```

```
In [47]: rfc=RandomForestClassifier()
model3 = rfc.fit(X_train, y_train)
model3.predict(X_test)
```

```
cm3=confusion_matrix(y_test, prediction3)
cm3
```

```
Out[47]: array([[109,  6],
                [ 6, 50]])
```

```
In [48]: sns.heatmap(cm3,annot=True, cmap ="BuPu",linecolor='White', linewidths=1)
plt.savefig('RandomForest.png')
```



```
In [49]: accuracy_score(y_test, prediction3)
```

```
Out[49]: 0.9298245614035088
```

```
In [50]: print(classification_report(y_test, prediction3))
```


	precision	recall	f1-score	support
0	0.95	0.95	0.95	115
1	0.89	0.89	0.89	56
accuracy			0.93	171
macro avg	0.92	0.92	0.92	171
weighted avg	0.93	0.93	0.93	171

```
In [51]: from sklearn import svm
```

```
In [52]: model4 = svm.SVC(kernel='rbf',C=30,gamma='auto')
model4.fit(X_train,y_train)
prediction4 = model4.predict(X_test)
cm4=confusion_matrix(y_test, prediction4)
cm4
```

```
Out[52]: array([[110,  5],
                [ 7, 49]])
```

```
In [53]: sns.heatmap(cm4,annot=True, cmap ="GnBu",linecolor='White', linewidths=1)
plt.savefig('SVM.png')
```



```
In [54]: accuracy_score(y_test, prediction4)
```

```
Out[54]: 0.9298245614035088
```

```
In [55]: print(classification_report(y_test, prediction4))
```


	precision	recall	f1-score	support
0	0.94	0.96	0.95	115
1	0.91	0.88	0.89	56
accuracy			0.93	171
macro avg	0.92	0.92	0.92	171
weighted avg	0.93	0.93	0.93	171

```
In [56]: from sklearn.naive_bayes import GaussianNB
```

```
In [57]: model5 = GaussianNB()
model5.fit(X_train,y_train)
prediction5 = model5.predict(X_test)
cm5=confusion_matrix(y_test, prediction5)
cm5
```

```
Out[57]: array([[110,  5],
[  7,  49]])
```

```
In [58]: sns.heatmap(cm5,annot=True, linecolor='white', linewidths=1)
plt.savefig('NaiveBayes.png')
```



```
In [59]: accuracy_score(y_test, prediction5)
```

```
Out[59]: 0.9298245614035088
```

```
In [60]: print(classification_report(y_test, prediction5))
```

	precision	recall	f1-score	support
0	0.94	0.96	0.95	115
1	0.91	0.88	0.89	56
accuracy			0.93	171
macro avg	0.92	0.92	0.92	171
weighted avg	0.93	0.93	0.93	171

```
In [61]: model_params = {
    'SVM': {
        'model': svm.SVC(gamma='auto'),
        'params': {
            'C': [1, 10, 20, 30, 40],
            'kernel': ['rbf', 'linear']
        }
    },
    'Random_Forest': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [1, 5, 10, 15, 20]
        }
    },
    'Logistic_Regression': {
        'model': LogisticRegression(solver='liblinear', multi_class='auto'),
        'params': {
            'C': [1, 5, 10, 15, 20]
        }
    },
    'Naive_Bayes_Gaussian': {
        'model': GaussianNB(),
```

```

        'params': {}
    },
    'Decision_Tree': {
        'model': DecisionTreeClassifier(),
        'params': {
            'criterion': ['gini', 'entropy'],
        }
    }
}

```

```

In [62]: from sklearn.model_selection import GridSearchCV
scores = []

for model_name, mp in model_params.items():
    clf = GridSearchCV(mp['model'], mp['params'], cv=5, return_train_score=False)
    clf.fit(X_train, y_train)
    scores.append({
        'model': model_name,
        'best_score': clf.best_score_,
        'best_params': clf.best_params_
    })

df1 = pd.DataFrame(scores, columns=['model', 'best_score', 'best_params'])
df1

```

```

Out[62]:

```

	model	best_score	best_params
0	SVM	0.927215	{'C': 10, 'kernel': 'rbf'}
1	Random_Forest	0.929684	{'n_estimators': 15}
2	Logistic_Regression	0.924620	{'C': 1}
3	Naive_Bayes_Gaussian	0.902120	{}
4	Decision_Tree	0.907025	{'criterion': 'entropy'}

In []: